

JIUCOM

MSA Infrastructure Manual

Docker + Monitoring + API Gateway + Distributed Tracing

9-Service Architecture Guide

Document : MSA Infrastructure Manual v1.0

Date : 2026-02-27

Project : JIUCOM API Platform

Stack : Spring Boot + Docker + Prometheus + Grafana

Table of Contents

1. Architecture Overview	9-Service MSA
2. Service Details	Each Container Role
3. Nginx API Gateway	Reverse Proxy + Load Balancing
4. Monitoring Stack	Prometheus + Grafana
5. Distributed Tracing	Zipkin + Micrometer
6. Docker Compose	Orchestration + Resource Limits
7. Data Flow	Request Lifecycle
8. Dashboard & Alerts	Grafana Panels
9. Operations Guide	Commands + Troubleshooting

1. Architecture Overview


>> Overall Architecture

JIUCOM uses a microservice-oriented Docker Compose architecture with 9 services. The system is designed around the following principles:

- 1) API Gateway Pattern - Nginx serves as the single entry point, routing requests to backend services and handling cross-cutting concerns (rate limiting, CORS, security headers).
- 2) Observability Stack - Prometheus collects metrics from all services, Grafana provides visualization dashboards, and Zipkin enables distributed tracing.
- 3) Sidecar Exporters - Redis Exporter and MySQL Exporter run as sidecar containers, translating native database metrics into Prometheus-compatible format.
- 4) Resource Isolation - Every container has CPU/memory limits to prevent resource starvation and simulate production-grade infrastructure.

>> Service Topology

[Client Browser]



2. Service Details

>> All 9 Services

Service	Image	Port	Role	Resources
mysql	mysql:8.0	3306	Primary database (user, part, build, post...)	2 CPU / 2GB
redis	redis:7-alpine	6379	Cache + session store (price cache, rate limit)	1 CPU / 512MB
app	Spring Boot (custom build)	8080	JIUCOM API server (70+ REST endpoints)	2 CPU / 1GB
nginx	nginx:alpine	80	API Gateway reverse proxy	0.5 CPU / 256MB
prometheus	prom/prometheus	9090	Metrics collection time-series DB	0.5 CPU / 512MB
grafana	grafana/grafana	3000	Monitoring dashboard visualization	0.5 CPU / 512MB
redis-exp	oliver006/ redis_exporter	9121	Redis metrics exporter	0.25 CPU / 128MB
mysql-exp	prom/ mysqld-exporter	9104	MySQL metrics exporter	0.25 CPU / 128MB
zipkin	openzipkin/ zipkin	9411	Distributed tracing request tracking	0.5 CPU / 512MB

>> Service Dependency Chain

Services start in dependency order:

- 1) mysql, redis (independent, start first with health checks)
- 2) app (depends on mysql + redis being healthy)
- 3) prometheus, zipkin (depend on app)

- 4) grafana (depends on prometheus)
- 5) redis-exporter, mysql-exporter (depend on redis/mysql)
- 6) nginx (depends on app, grafana, zipkin - starts last)

3. Nginx API Gateway

>> Role & Responsibilities

Nginx acts as the API Gateway - the single entry point for all external traffic. This is a core MSA pattern that provides:

- * Reverse Proxy - Routes requests to the correct backend service based on URL path
- * Rate Limiting - Protects backend from DDoS/abuse (30 req/sec per IP, burst 50)
- * Security Headers - X-Frame-Options, X-Content-Type-Options, XSS protection
- * WebSocket Support - HTTP/1.1 Upgrade for real-time notifications (STOMP)
- * Gzip Compression - Reduces response payload size for JSON/text

>> Routing Rules

Path	Upstream	Description
/api/v1/*	app:8080	All REST API endpoints (auth, parts, builds, posts...)
/grafana/*	grafana:3000	Monitoring dashboard (Grafana UI)
/zipkin/*	zipkin:9411	Distributed tracing UI (Zipkin web console)
/health	(direct)	Nginx health check returns 200 OK
/	(redirect)	Redirects to /api/v1/swagger-ui.html

>> How Reverse Proxy Works

When a client sends a request to <http://localhost/api/v1/part>:

- 1) Nginx receives the request on port 80
- 2) The 'location /api/v1/' block matches the URL
- 3) Rate limiter checks if the client IP is within the 30 req/sec limit

- 4) If allowed, Nginx forwards (proxies) the request to `http://app:8080/api/v1/parts`
- 5) Nginx adds headers: X-Real-IP, X-Forwarded-For, X-Forwarded-Proto
- 6) The Spring Boot app processes the request and returns a response
- 7) Nginx adds security headers and forwards the response to the client

The client never communicates directly with the Spring Boot app. This provides a layer of abstraction and security.

4. Monitoring Stack

>> Prometheus - Metrics Collection

Prometheus is a time-series database that collects metrics from all services using a pull-based model. Every 15 seconds, Prometheus scrapes (HTTP GET) each target's metrics endpoint.

Prometheus Pull Model: Prometheus --HTTP GET--> Target /metrics

Scrape targets configured in prometheus.yml:

Job Name	Target	Metrics Path	What It Collects
jiucom-api	app:8080	/api/v1/actuator/	JVM heap, GC, threads
		prometheus	HTTP requests, HikariCP
redis-exporter	redis-exporter :9121	/metrics	Connected clients
			memory usage, commands
mysql-exporter	mysql-exporter :9104	/metrics	Queries/sec, threads
			slow queries, connections

>> How Spring Boot Exposes Metrics

The Spring Boot app uses Micrometer (a metrics facade) to expose metrics:

- 1) spring-boot-starter-actuator enables the /actuator endpoints
- 2) micrometer-registry-prometheus formats metrics in Prometheus text format
- 3) application.yml exposes: health, info, metrics, prometheus, loggers
- 4) Prometheus scrapes GET /api/v1/actuator/prometheus every 15s

Key metrics exposed:

- * jvm_memory_used_bytes - JVM heap/non-heap memory usage
- * jvm_threads_live_threads - Number of active threads
- * http_server_requests_seconds_count - HTTP request count by status/method/uri
- * http_server_requests_seconds_sum - Total response time (for avg calculation)
- * hikaricp_connections_active - Active database connections

* jvm_gc_pause_seconds_count - Garbage collection frequency

>> Grafana - Visualization Dashboard

Grafana reads data from Prometheus and displays it as visual dashboards. The JIUCOM dashboard is auto-provisioned on startup.

Auto-provisioning files (mounted as Docker volumes):

- * datasource.yml - Connects Grafana to Prometheus (<http://prometheus:9090>)
- * dashboard.yml - Tells Grafana to load dashboards from /var/lib/grafana/dashboards/
- * jiucom-dashboard.json - Pre-built dashboard with 16 panels

>> Dashboard Panels (16 Total)

Section	Panels	Metrics Used
Application	API Status, Heap Usage	up{}, jvm_memory_*
Overview	Threads, Request Rate	jvm_threads_*
	Error Rate (5xx)	http_server_requests_*
HTTP	Response Time (avg)	http_server_requests_
	Requests by Status	seconds_sum/count
JVM	Heap Memory	jvm_memory_used_bytes
	(used/committed/max)	jvm_gc_pause_*
Database & Cache	GC Pause Count	
	HikariCP Pool	hikaricp_connections_*
	Redis Clients	redis_connected_clients
MySQL	Redis Memory	redis_memory_used_bytes
	Queries/sec	mysql_global_status_
	Slow Queries	queries/threads_*
	Threads	

>> Exporter Pattern

Redis and MySQL don't natively expose Prometheus metrics. Exporters are sidecar containers that:

- 1) Connect to the database (Redis/MySQL) using native protocol
- 2) Query internal statistics (INFO, SHOW GLOBAL STATUS)

- 3) Translate those stats into Prometheus metric format
- 4) Expose them on /metrics endpoint for Prometheus to scrape

This is the standard Prometheus 'Exporter' pattern for third-party systems.

5. Distributed Tracing (Zipkin)

>> What is Distributed Tracing?

In a microservice architecture, a single user request may pass through multiple services. Distributed tracing tracks the entire journey of a request across all services, showing:

- Which services were called
- How long each service took
- Where bottlenecks or errors occurred

Each request gets a unique Trace ID, and each service call within that request gets a Span ID.

>> How It Works in JIUCOM

- 1) Client sends request to Nginx -> forwarded to Spring Boot app
- 2) Micrometer Tracing (Brave bridge) automatically:
 - Generates a Trace ID for the request
 - Creates a Span for each operation (HTTP handler, DB query, Redis call)
 - Propagates trace context via HTTP headers (B3 propagation)
- 3) When the request completes, spans are sent asynchronously to Zipkin
via HTTP POST to `http://zipkin:9411/api/v2/spans`
- 4) Zipkin stores the trace data (in-memory for dev, could use Elasticsearch for prod)
- 5) Zipkin UI (port 9411) visualizes the trace as a timeline/waterfall diagram

>> Dependencies Added

```
// build.gradle.kts  
implementation("io.micrometer:micrometer-tracing-brave")  
implementation("io.zipkin.reporter2:zipkin-reporter-brave")  
  
// application.yml  
management.tracing.sampling.probability: 1.0 (dev: 100%)  
management.zipkin.tracing.endpoint: http://zipkin:9411/api/v2/spans
```

>> Sampling Strategy

Profile	Sampling Rate	Reason
dev	1.0 (100%)	Trace every request for debugging
prod	0.1 (10%)	Reduce overhead, sample enough for analysis

6. Docker Compose Orchestration

>> Container Resource Limits

Every container has CPU and memory limits to prevent resource starvation. This simulates production-grade Kubernetes resource management.

Service	CPU Limit	Memory Limit	CPU Reserve	Memory Reserve
app	2.0	1 GB	0.5	512 MB
mysql	2.0	2 GB	0.5	512 MB
redis	1.0	512 MB	0.25	128 MB
nginx	0.5	256 MB	0.1	64 MB
prometheus	0.5	512 MB	0.25	256 MB
grafana	0.5	512 MB	0.25	256 MB
zipkin	0.5	512 MB	0.25	256 MB
redis-exp	0.25	128 MB	0.1	64 MB
mysql-exp	0.25	128 MB	0.1	64 MB

>> Network Architecture

All services share a single Docker bridge network: jiucom-network

Within this network, containers communicate using service names as DNS hostnames. For example, the app container connects to 'mysql:3306' and 'redis:6379'. Docker's embedded DNS resolves these names to container IP addresses.

Only specific ports are published to the host machine:

- Port 80 (Nginx) - main entry point
- Port 8080 (App) - direct API access (for development)
- Port 3000 (Grafana), 9090 (Prometheus), 9411 (Zipkin) - monitoring

>> Volume Persistence

Volume	Container Path	Purpose
--------	----------------	---------

mysql-data	/var/lib/mysql	MySQL data files (survives container restart)
redis-data	/data	Redis RDB snapshots
prometheus-data	/prometheus	Time-series metrics data (15 day retention)
grafana-data	/var/lib/grafana	Grafana settings, users (not dashboards - those are provisioned)

7. Data Flow - Request Lifecycle

>> Example: GET /api/v1/parts (Search Parts)

Step-by-step flow when a client searches for computer parts:

Phase 1: Request Ingress

- * Client sends HTTP GET http://localhost/api/v1/parts?category=CPU
- * Nginx receives on port 80, matches 'location /api/v1/'
- * Rate limiter checks: client IP within 30 req/sec? If yes, proceed
- * Nginx proxies to http://app:8080/api/v1/parts?category=CPU
- * Adds headers: X-Real-IP, X-Forwarded-For, X-Forwarded-Proto

Phase 2: Application Processing

- * Spring Boot receives request, Micrometer creates Trace ID + Span
- * JwtAuthenticationFilter checks Authorization header (if present)
- * RequestLoggingFilter logs: method, URI, client IP, start time
- * PartController.searchParts() invoked by DispatcherServlet
- * PartService calls PartRepository (QueryDSL dynamic query)
- * HikariCP provides a connection from the pool to MySQL
- * Query executes, results mapped to PartSearchResponse DTOs

Phase 3: Response Egress


- * ApiResponse<T> wraps the result: {success: true, data: [...], message: ...}
- * Micrometer records: http_server_requests_seconds (duration, status, uri)
- * Trace span sent async to Zipkin (http://zipkin:9411/api/v2/spans)
- * Response returns through Nginx, security headers added
- * Client receives JSON response with parts data

Phase 4: Metrics Collection (Background)

- * Every 15 seconds, Prometheus scrapes GET /api/v1/actuator/prometheus
- * Metrics collected: request count, response time, JVM heap, thread count
- * Prometheus stores in TSDB (time-series database), 15-day retention
- * Grafana queries Prometheus via PromQL for dashboard visualization
- * Redis Exporter scrapes Redis INFO, MySQL Exporter scrapes SHOW STATUS

>> Monitoring Data Flow

Spring Boot App ----metrics----> Prometheus ----query----> Grafana



8. Dashboard & Alert Design

>> Grafana Dashboard Layout

The pre-provisioned JIUCOM Dashboard (UID: jiucom-main) has 4 row sections with 16 panels total. Auto-refresh is set to every 10 seconds.

Row	Panel Type	What to Watch	Warning Sign
Overview (Row 1)	Stat panels x5	API UP status	DOWN status
		Heap usage %	Heap > 90%
		HTTP request rate	Rate spike/drop
HTTP (Row 2)	Time series x2	Avg response time	Latency > 1s
		Requests by status code	5xx errors > 1%
JVM (Row 3)	Time series x2	Heap memory trend	Memory climbing
		GC pause frequency	Frequent full GC
DB/Cache (Row 4)	Time series x5	HikariCP pool usage	Pool exhaustion
		Redis clients/memory	Redis memory full
		MySQL queries/threads	Slow query spike

>> Key PromQL Queries

```
# API Server Up/Down
up{job="jiucom-api"}
```

```
# JVM Heap Usage Percentage
jvm_memory_used_bytes{area="heap"} / jvm_memory_max_bytes{area="heap"} * 100
```

```
# HTTP Request Rate (per second, 5m average)
rate(http_server_requests_seconds_count{job="jiucom-api"}[5m])
```

```
# HTTP 5xx Error Rate (%)
rate(http_server_requests_seconds_count{status=~"5.."}[5m])
/ rate(http_server_requests_seconds_count[5m]) * 100
```

```
# Average Response Time (seconds)
rate(http_server_requests_seconds_sum[5m])
/ rate(http_server_requests_seconds_count[5m])
```

9. Operations Guide

>> Start / Stop Commands

```
# Start all 9 services (build app image first)
docker compose up -d --build
```

```
# Start without rebuilding
docker compose up -d
```

```
# Stop all services (keep data volumes)
docker compose down
```

```
# Stop and delete all data (DESTRUCTIVE)
docker compose down -v
```

```
# View logs for specific service
docker compose logs -f app
docker compose logs -f nginx
```

```
# Restart single service
docker compose restart app
```

>> Health Check URLs

Service	URL	Expected
API (via Nginx)	http://localhost/api/v1/	200 OK
	actuator/health	{status: UP}
API (direct)	http://localhost:8080/	200 OK
	api/v1/actuator/health	{status: UP}
Prometheus	http://localhost:9090/	3 targets UP
	targets	
Grafana	http://localhost:3000/	Login page
	grafana/login	(admin/admin)

Zipkin	http://localhost:9411/ zipkin/	Zipkin UI
Swagger	http://localhost/api/v1/ swagger-ui.html	Swagger UI

>> Troubleshooting

Problem	Diagnosis	Solution
App won't start	docker compose logs app	Ensure mysql is healthy
	Check DB connection	Check DB_HOST env var
Prometheus no targets	http://localhost:9090	Check prometheus.yml
	/targets	Ensure app is running
Grafana no data	Check datasource	Verify Prometheus URL
	settings in Grafana	http://prometheus:9090
Nginx 502 Bad Gateway	docker compose logs nginx	Ensure app container
	Check upstream	is running and healthy
Zipkin no traces	Check app logs for	Verify ZIPKIN_ENDPOINT
	Zipkin reporter errors	env var in app
MySQL exporter connection refused	docker compose logs	Check DATA_SOURCE_NAME
	mysql-exporter	MySQL must be healthy

>> Scaling Considerations

For production deployment, consider these upgrades:

- 1) Nginx: Add SSL/TLS termination with Let's Encrypt certificates
- 2) App: Scale horizontally with 'docker compose up --scale app=3'
- 3) Zipkin: Switch from in-memory to Elasticsearch storage
- 4) Prometheus: Add alertmanager for Slack/email notifications
- 5) Grafana: Configure SMTP for scheduled report delivery
- 6) MySQL: Add read replicas for query distribution
- 7) Redis: Enable Redis Cluster for high availability

JIUCOM - Production-Ready MSA Infrastructure

Document generated: 2026-02-27 | JIUCOM Team